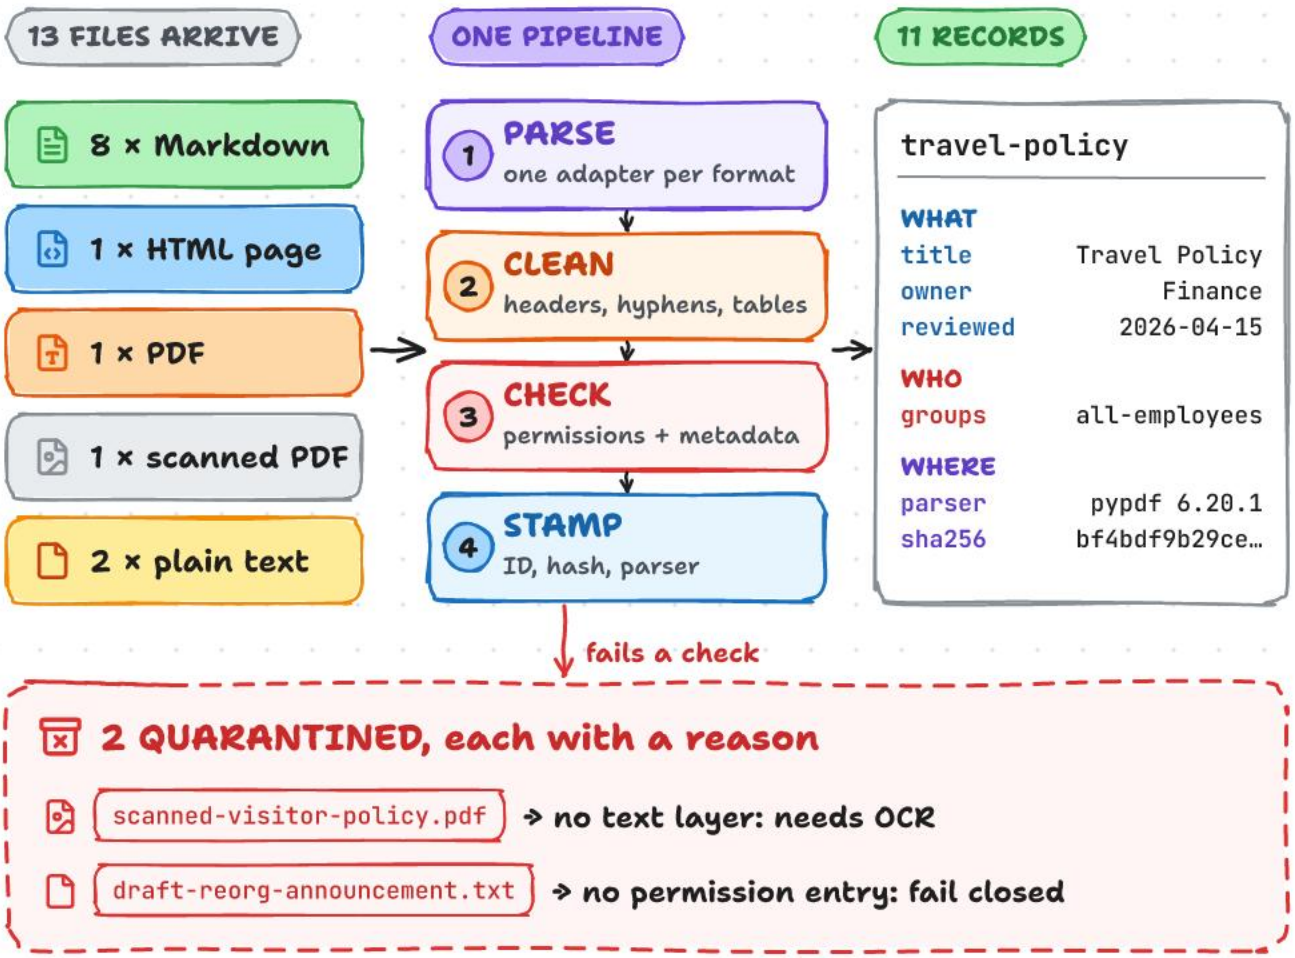
Messy files in, trusted records out

Day 02: ingestion, parsing and metadata, before any retrieval.



Retrieval can only find what ingestion kept.
So check what goes in, and say what was left out.

Real run · Day 02 notebook · free, no API key

Where each field comes from

Who can change a field decides how far you can trust it.

THE DOCUMENT ITSELF

front matter · <meta> tags · PDF properties

`title` `owner` `last_reviewed`

⚠ **Any author can edit it. Validate it.**

e.g. intranet page: review overdue (470 days)

THE SYSTEM OF RECORD

the permission export, not the file

`allowed_groups` **authoritative · missing = fail closed**

A dropped-in draft says:

"Audience: all staff (please share widely)"

The export says:

no entry →

visible to nobody

THE PIPELINE

added by your code on every run

`doc_id` `source_path` `parser` `content_hash`

Traces any answer to one file, one parser version and one exact version of the text.

PDFs store ink, not tables

One page of travel-policy.pdf, extracted two ways by pypdf 6.20.1.

WHAT YOU SEE

City	Hotel per night	Meals per day
London	USD 250	USD 75
New York	USD 300	USD 75
Singapore	USD 220	USD 60
Other cities	USD 180	USD 50

DEFAULT ("plain") MODE

follows the drawing order

City | London | New York | Singapore | Other cities | Hotel per night | USD 250 | USD 300 | USD 220 | USD 180 | Meals per day | USD 75 | USD 75 | USD 60 | USD 50

One stream: all cities, then all hotel limits, then all meal limits.

LAYOUT MODE + CLEANING

places text by position

City | Hotel per night | Meals per day

London | USD 250 | USD 75

New York | USD 300 | USD 75

Singapore | USD 220 | USD 60

Other cities | USD 180 | USD 50

Rows survive: Singapore → USD 220

claim reim- / bursement

re-joined → phrase found

header + footer on each page

4 lines → 0

Same document, new version

Two IDs, two questions. Compare them on every run.

WHICH document?

`doc_id` stays when the text changes

`remote-work-policy`

WHICH version?

`content_hash` changes with 1 character

USD 500 `9d0e7b74370d6210…`

USD 600 `e6e17cd1b3d6f8db…`

Real run: 5 changes to a copy of the data, then re-ingest

ADDED
`annual-leave-policy`
→ chunk + index

TEXT CHANGED
`remote-work-policy`
→ re-chunk + re-index

PERMISSIONS CHANGED
`password-and-mfa-standard +contractors`
→ update metadata only

REMOVED
`vpn-access-guide`
→ DELETE its chunks

UNCHANGED
`8 docs (one re-saved)`
→ skip

One file was re-saved with Windows line endings: new bytes, same text, so it is skipped.

Next → Day 03: chunking and document structure